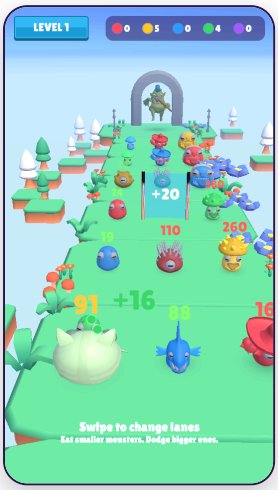
1 · THE AD

A 15–second storyboard

Portrait 9:16, readable with the sound off: the payoff first (a bite and a big number), then the loop that makes it, then the twist that sets the game apart, the painting.

0–2 s · Hook

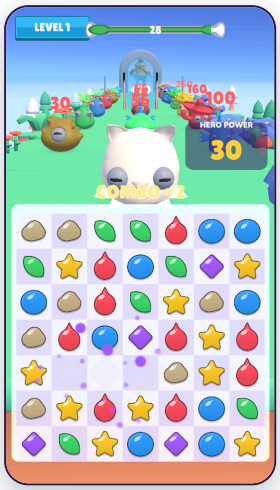


The hero bites a monster: +N pops, paint splashes.

EAT THE SMALL ONES

Sound: Chomp, gulp

2–4 s · The rule

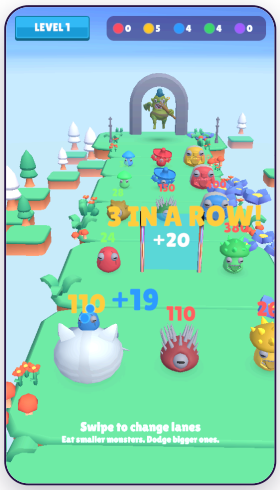


A cascade on the board; sparks fly into the hero, who swells.

MATCH TO GROW

Sound: Rising chimes

4–6.5 s · The run

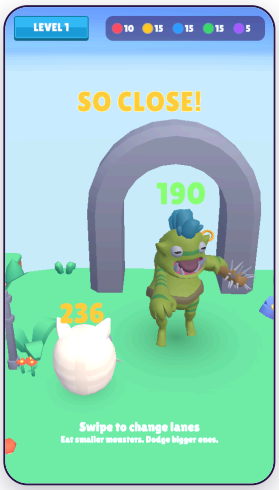


Swipe, bite, bite, bite; a +20 gate.

3 IN A ROW!

Sound: Chomps climbing in pitch

6.5–8.5 s · Tension

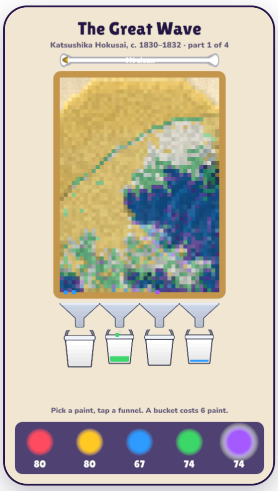


The guardian: slow motion, the hero wins by a hair.

236 vs 190

Sound: Roar, slow bite, fanfare

8.5–11 s · The twist



The paint becomes sand pouring off the Great Wave into buckets.

NOW SAVE THE MASTERPIECE

Sound: Sand ticks, bucket ding

11–12.5 s · Payoff

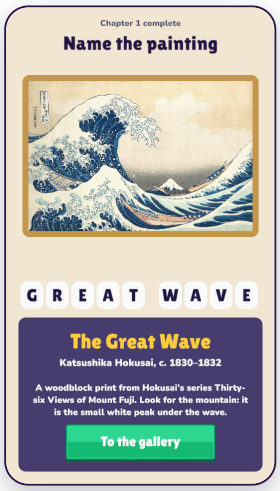


A brush sweeps the last sand; the painting appears.

—

Sound: Brush swish

12.5–14 s · Collect

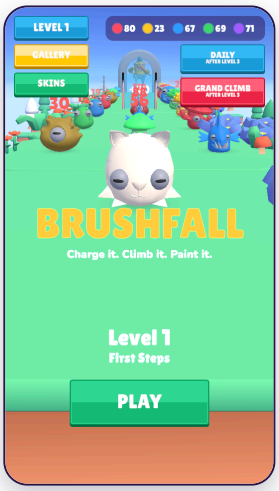


The Great Wave named, into the gallery.

RESTORE THE WORLD'S ART

Sound: Fanfare

14–15 s · End card



Title over the hero; Play pulses; store badges.

PLAY FREE

Sound: Music button

Every frame is real. A PlayMode test (`AdStoryboard_CapturesTheBeats`) plays these moments in the build and saves them, setting the charge to the power that beats the guardian by a hair. The ad's promise is the first minute of play, which is the opposite of the "fake" power-climb ads this audience already distrusts.

Three hooks, one decision

The prototype's first question is not "is it fun?" but "will people install it at a price we can pay back?"

Creative variants

VARIANT	FIRST 2 SECONDS	WHAT IT TESTS
A · Payoff (above)	A bite: the hero eats a smaller number, +N pops	The climb's satisfaction, the widest audience
B · Fail	The hero walks into a bigger number and is eaten. "Can you pick the right lane?"	The challenge hook runner ads use
C · Masterpiece	Sand pouring off the Great Wave into coloured buckets	The puzzle and art audience; is the restoration the stronger hook?

Test design

	SETTING
Channels	Meta (Facebook/Instagram) and one ad network (AppLovin or Unity Ads), where hybrid-casual tests usually run
Geos	US for the go / kill decision; a cheaper tier-2 geo (Brazil, Turkey) to iterate the creatives first
Build	The Android build behind a store page (or a soft-launch listing), analytics on
Budget	About $150–250 per variant in the cheap geo, then the winner at about $500 in the US
Length	3–5 days per round, at least 1,000 impressions per variant before reading CTR
Rounds	Up to three creative rounds before the kill decision

Go / iterate / kill (US Android, Meta)

Go

- CPI ≤ $0.50
- CTR ≥ 3%; IPM ≥ 40 on the network
- D1 ≥ 35% from the test cohort
- Day-0 playtime ≥ 15 minutes

Soft launch with the KPI plan.

Iterate

- CPI $0.50–0.90, or
- good CTR but weak store conversion, or
- D1 25–35%

Change the hook first (cheapest), then the store page (icon, first screenshot), then the first minute of play.

Kill

- CPI above $0.90 after three creative rounds, or
- D1 below 25% even with a good CPI

The fantasy is not landing: test a new theme (the climb with a different reward) before more content.

Why these numbers: hybrid-casual games earn from ads and some IAP, so a practical target is day-30 revenue per user above the install price. With the KPI targets overleaf (D1 40%, D7 15%, ad ARPDAU ≈ $0.07) a user earns roughly $0.45–0.65 by day 30, which is why the go line sits at a $0.50 CPI. These are planning assumptions, to be replaced by the publisher's own benchmarks.

3 · KPI PLAN

What we measure, and what we change

Every KPI maps to an event the game already logs through `IAnalytics` (in this build, a local JSON-lines file per device; a real SDK drops in behind the same interface).

KPI	TARGET	MEASURED BY (DOCS/ANALYTICS.MD)	IF IT MISSES
D1 retention	40%	`session_start` on day 1 per install	Onboarding: level 1–3 funnel, first restoration
D7 retention	15%	`session_start` on day 7	Chapter pacing, hard levels, the daily habit
D30 retention	5%	`session_start` on day 30	Content depth: chapters, the Grand Climb
Daily painting played	40% of DAU	`daily_start` / sessions	Button placement, reward value, streak
Grand Climb reach	50% of weekend DAU	`grand_end` per weekend	Unlock timing, tier rewards
First painting named	70% of installs	`finale_end` chapter 1	Restoration difficulty, paint shortfall
Restoration out of paint	≤ 25% of visits	`restoration_end` OutOfPaint	Bucket price, colour drops (workbook)
Rewarded per DAU	1.5–2.5	`ad_rewarded_end` earned	Offer value; too many offers is fatigue
Ad ARPDAU	$0.05–0.09	ad SDK	Placement mix, interstitial pacing
Payer conversion	1–2%	`purchase_end` Success	Starter pack value, shop visibility

First-session funnel (targets)



Read difficulty in the wild

`climb_end` (won, row reached, revived) by level is the real version of the bots' win-rate chart. Where players and bots disagree, the bots are wrong: retune the level and rerun the sim.

Read the economy

`restoration_end` with OutOfPaint followed by an extra climb is the paint shortfall in the wild; wallet totals at `session_start` and `skin_buy` show whether the sink keeps up.